

Job Match Report

Search keywords: AI Engineer RAG LLM multi-agent

Experience level: associate

Generated: 2026-04-28 10:55:12

Mobin Azimipannah is an AI Engineer with dual M.Sc. degrees in Electrical Engineering and Computer Science, specializing in machine learning, NLP, and retrieval-augmented generation (RAG). Currently at Hermes Corporate, he consults on AI infrastructure, having previously developed an automated RAG system at Canon Research Center, fine-tuned LLMs using LoRA/quantization, and built multi-agent applications. His master's thesis explored a Multi-RAG Framework, and he led a multi-agent LLM project at a Berkeley hackathon. Skilled in PyTorch, TensorFlow, LangChain, AutoGen, FastAPI, and AWS, with a strong full-stack foundation. EU work authorized.

Netherlands

1. AI Engineer — Reply

- **Location:** Milan, Lombardy, Italy
- **Match score:** 100/100
- **Apply:** <https://www.linkedin.com/jobs/view/4402863479/>

Why it fits: The candidate holds a master's from Politecnico di Milano (where the job is located) and currently works as an AI Engineer with strong experience in LLMs, RAG, agents, and full-stack AI platforms, directly matching the job title. No job description is provided, but based on the title and location, all evidence points to a perfect fit with no apparent gaps.

Cover letter

Dear Hiring Team,

I am writing to apply for the AI Engineer position at Reply in Milan. I hold EU work authorization and require no visa sponsorship. Currently based in Italy, I will soon complete an M.Sc. in Electrical Engineering – Computer Science at Politecnico di Milano, where my thesis focuses on a Multi-RAG Framework with Prof. Mark James Carman. My academic background is complemented by hands-on engineering experience and a strong interest in building production-ready AI systems.

In my current role as an AI Engineer at Hermes Corporate in Florence, I consult for Baker Hughes on AI infrastructure. I develop OpenAPI and MCP services for surrogate model training, DOE generation, preprocessing, and reporting. The stack includes FastAPI, PydanticAI, PostgreSQL, MLflow, Docker, and NGINX, all deployed on AWS. This experience has given me solid exposure to end-to-end MLOps and API design in industrial settings.

Previously, at Canon Research Centre France, I led the development of an automated RAG system that reduced IP document matching from several days to under three minutes while maintaining over 90% coverage. I designed hierarchical ensemble retrievers and integrated AutoGen with GraphRAG for multi-agent workflows. I also fine-tuned local LLMs using LoRA, MoRA, VeRA, and quantization methods.

I complement my work with hands-on projects. I am building Zirak.dev, a full-stack AI assistant platform with Next.js, React, and a custom Python agentic workflow that autonomously decomposes complex tasks. During the 2024 UC Berkeley LLM Agent Hackathon, I led the architecture of regulAItor, a multi-agent system using AutoGen to analyze over 1,000 FDA warning letters. Additionally, I engineered a real-time cryptocurrency prediction system with Kafka, InfluxDB, and xLSTM models deployed via Docker.

My technical stack spans PyTorch, TensorFlow, LangChain, LlamaIndex, AutoGen, GraphRAG, FastAPI, Docker, and AWS. I also hold multiple certifications in LLMs, NLP, computer vision, and machine learning from DeepLearning.AI, Stanford, and Google Cloud.

I am eager to bring my background in AI engineering, system design, and multi-agent orchestration to Reply's projects in Milan. Thank you for considering my application. I look forward to the opportunity to discuss how I can contribute to your team.

Sincerely, Mobin Azimipناه

2. AI Engineer — TextYess

- **Location:** Remote
- **Match score:** 95/100
- **Apply:** <https://www.linkedin.com/jobs/view/4401625493/>

Why it fits: The candidate holds an M.Sc. in AI/ML and works as an AI Engineer with recent experience building AI infrastructure (FastAPI, AWS, Docker). Strong project portfolio includes advanced RAG, LLM fine-tuning, and full-stack AI apps, closely matching typical AI Engineer responsibilities. Main gap is lack of job description to confirm domain-specific requirements, but generic fit is excellent.

Cover letter

Subject: Application for AI Engineer at TextYess

Dear Hiring Team,

I am writing to express my interest in the AI Engineer position at TextYess. My background spans building production AI infrastructure, developing retrieval-augmented generation (RAG) systems, and engineering LLM-powered applications end-to-end—all of which I believe can bring immediate value to your team in a remote capacity.

I am currently an AI Engineer at Hermes Corporate, consulting for Baker Hughes. I build OpenAPI and MCP services for surrogate model training, design-of-experiments generation, and automated reporting. My daily work involves FastAPI, PydanticAI, PostgreSQL, MLflow, Docker, and NGINX, all deployed on AWS. This role requires shipping scalable, maintainable code and collaborating closely with domain experts, habits I apply to every AI project.

During my internship at Canon Research Center France, I led development of an automated RAG system that matched intellectual property documents with over 90% coverage in under three minutes—a task that previously took researchers days. I designed hierarchical ensemble retrievers, integrated AutoGen and GraphRAG for multi-agent workflows, and fine-tuned local LLMs using LoRA, MoRA, VeRA, and quantization. That project gave me hands-on expertise in retrieval pipelines, evaluation, and adapting large models to specialized tasks.

My academic work reinforces these skills. I am completing an M.Sc. at Politecnico di Milano

with a thesis on a Multi-RAG framework. I also built a multi-agent LLM application for pharmaceutical regulatory compliance at the 2024 UC Berkeley LLM Agent Hackathon, and I engineered a full-stack AI assistant platform using Next.js, React, and a custom Python agentic workflow. These experiences have made me fluent in Python, PyTorch, Transformers, LangChain, AutoGen, Docker, and AWS.

I have delivered remote freelance prompt-engineering strategies for frontier model training through Outlier, so I am comfortable working independently and communicating effectively across time zones. I am eager to contribute my blend of research depth and production engineering to TextYess.

Thank you for considering my application.

Sincerely,
Mobin Azimipannah
+39 3516515966
azimipannah.mobin@gmail.com

3. AI Engineer — Pillar

- **Location:** Milan, Lombardy, Italy
- **Match score:** 95/100
- **Apply:** <https://www.linkedin.com/jobs/view/4401964351/>

Why it fits: The candidate's strong background in AI engineering, including LLMs, RAG, agentic workflows, and full-stack AI platforms, aligns perfectly with a typical AI Engineer role. Their education at Politecnico di Milano and current presence in Italy, along with EU work authorization, make them an ideal location match. The only minor gap is the absence of a detailed job description to confirm specific requirements, but the CV demonstrates all the expected competencies.

Cover letter

Dear Hiring Team at Pillar,

I am writing to express my interest in the AI Engineer position at Pillar in Milan. As a Master's candidate in Electrical Engineering and Computer Science at the Polytechnic University of Milan, with dual specialization in machine learning and robotics, I bring directly relevant AI engineering experience, full proficiency in the European job market, and a strong connection to Milan.

Currently, I work as a full-time AI Engineer at Hermes Corporate in Florence, consulting Baker Hughes on AI infrastructure. I design and deploy services for surrogate model training, DOE generation, preprocessing, and reporting using FastAPI, PydanticAI, PostgreSQL, MLflow, Docker, and NGINX, all on AWS. This role has sharpened my ability to deliver production-grade AI systems end-to-end. Prior to this, I led the development of an automated RAG system during my internship at Canon Research Center France. The system achieved over 90% coverage of intellectual property content and references in under three minutes—a task that previously took researchers days. I built hierarchical ensemble retrievers and integrated AutoGen with GraphRAG to coordinate multi-agent workflows, and I fine-tuned local LLMs with LoRA, MoRA, VeRA, and quantization.

My project portfolio reinforces this production mindset. I engineered Zirak.dev, a full-stack AI assistant platform with a custom Python-based agentic workflow that autonomously

decomposes complex tasks and manages completion loops, built with Next.js and React. I also led the technical architecture of regulAItor, a multi-agent LLM system using AutoGen to streamline pharmaceutical regulatory compliance, which won recognition at the UC Berkeley LLM Agent Hackathon. In addition, I have delivered over 50 innovative prompt engineering strategies for frontier model training as a freelance AI Expert at Outlier.

Academically, my thesis (Multi-RAG Framework, supervised by Prof. Mark James Carman) and research on sparse techniques in neural network backpropagation under Prof. Agosta at Politecnico di Milano have deepened my understanding of efficient, scalable AI. I am proficient across the modern AI stack: PyTorch, TensorFlow, Transformers, LangChain, Llama-index, AutoGen, MCP, MLflow, Docker, AWS, and both backend (Python, FastAPI, PostgreSQL) and frontend (Next.js, React) technologies.

I hold EU work authorization and require no visa sponsorship. I am eager to bring my blend of hands-on AI engineering, agentic systems, and full-stack development to Pillar's team in Milan. I would welcome the opportunity to discuss how my background aligns with your needs. Thank you for your consideration.

Sincerely,
Mobin Azimipannah

4. Agentic AI Expert — Capgemini

- **Location:** Milan, Lombardy, Italy
- **Match score:** 92/100
- **Apply:** <https://www.linkedin.com/jobs/view/4401910375/>

Why it fits: The candidate demonstrates strong agentic AI expertise through multiple projects and work experience directly using AutoGen, multi-agent systems, GraphRAG, and RAG frameworks. Their education includes a thesis on Multi-RAG and they are based in Italy with EU work authorization, perfectly aligning with the Milan location. The main gap is the lack of a detailed job description to assess specific industry or seniority requirements, but the overall match is excellent.

Cover letter

Dear Hiring Manager,

I am writing to express my strong interest in the Agentic AI Expert position at Capgemini in Milan. With a Master's degree from Politecnico di Milano, a thesis on a Multi-RAG Framework, and extensive hands-on experience building multi-agent LLM systems, I am well-prepared to design and implement sophisticated agentic solutions.

During my internship at Canon Research Center France, I led the development of an automated RAG system that combined hierarchical ensemble retrievers with AutoGen and GraphRAG to orchestrate multi-agent workflows. This system analyzed intellectual property documents and their references, achieving over 90% coverage in under three minutes—a task that previously took researchers days. That work solidified my ability to leverage multi-agent architectures for real-world, knowledge-intensive tasks.

At the 2024 UC Berkeley Large Language Model Agent Hackathon, I served as technical architect for regulAItor, a multi-agent LLM application built on AutoGen that analyzed over 1,000 FDA warning letters to streamline pharmaceutical regulatory compliance. I am currently extending these concepts through Zirak.dev, a full-stack AI assistant platform

where I engineered a custom Python-based agentic workflow that autonomously decomposes complex multi-step tasks and manages completion loops.

In my current role as AI Engineer at Hermes Corporate, I consult for Baker Hughes on AI infrastructure, building OpenAPI and MCP services for surrogate model training and deployment with FastAPI, PydanticAI, MLflow, Docker, and AWS. Additionally, I deliver advanced prompt engineering strategies as a freelancer for Outlier, helping train frontier models. This blend of agentic R&D and production-grade engineering gives me a pragmatic perspective on deploying reliable AI agents at scale.

I hold European work authorization and have been based in Italy for the past three years, splitting my time between Florence and Milan. I am eager to bring my expertise in multi-agent LLMs, RAG, and robust AI infrastructure to Capgemini's projects in Milan. Thank you for considering my application.

Sincerely,
Mobin Azimipannah

5. Software Engineer - AI — Jet HR

- **Location:** Remote
- **Match score:** 90/100
- **Apply:** <https://www.linkedin.com/jobs/view/4401922191/>

Why it fits: Strong match due to relevant AI engineering role, advanced degree in ML/Robotics, and extensive project portfolio in LLMs, RAG, and full-stack AI. Main gaps include early career stage and lack of job description to confirm specific tool alignment, but overall skills are highly transferable.

Cover letter

Dear Hiring Team at Jet HR,

I am writing to express my interest in the Software Engineer - AI position at Jet HR. With a strong background in building AI-powered systems, full-stack development, and cloud deployment, I am eager to apply my skills to help Jet HR create impactful HR technology solutions.

Currently, I work as an AI Engineer at Hermes Corporate, where I consult for Baker Hughes on AI infrastructure. I design and deploy services using FastAPI, PydanticAI, PostgreSQL, MLflow, Docker, NGINX, and AWS, including OpenAPI and MCP-based surrogate model training and DOE generation pipelines. This experience has sharpened my ability to deliver production-grade AI systems.

During my internship at Canon Research Center France, I led the development of an automated RAG system that identified relevant matches between intellectual property documents. I used hierarchical ensemble retrievers, AutoGen with GraphRAG, and multi-agent workflows to achieve over 90% coverage in under three minutes—work that previously took researchers days. I also fine-tuned local LLMs using LoRA, MoRA, VeRA, and quantization methods. These projects gave me deep expertise in retrieval-augmented generation, multi-agent coordination, and efficient model optimization.

My freelance work as an AI Expert and Prompt Engineer has involved delivering over 50 prompt engineering strategies to train frontier models, further strengthening my understanding of large language model behavior. Additionally, I have built full-stack AI

platforms like Zirak.dev using Next.js, React, and a custom Python agentic workflow that autonomously decomposes complex tasks. At the UC Berkeley LLM Agent Hackathon, I architected a multi-agent system leveraging AutoGen to analyze FDA warning letters, streamlining regulatory compliance. I also contributed to a multi-LLM chatbot for pet care that combined AutoGen, RAG, and MongoDB.

My academic background includes a Master's in Electrical Engineering – Computer Science from Polytechnic University of Milan, where my thesis focused on a Multi-RAG Framework. Throughout my studies and projects, I have consistently worked at the intersection of machine learning, full-stack engineering, and cloud deployment, ready to tackle the challenges of building intelligent, scalable HR tools.

I am excited by the opportunity to bring my cross-functional AI engineering experience to Jet HR and contribute to solutions that can transform HR workflows. Thank you for considering my application. I look forward to the possibility of discussing how I can add value to your team.

Best regards, Mobin Azimipناه +39 3516515966 azimipناه@gmail.com